

CAS in Applied Data Science


Muerren 2026



Géraldine Schaller-Conti

Bibliography


- Deep Learning book (Goodfellow, Bengio, Courville)
- Machine Learning @ Stanford (Prof Andrew Ng)
- Hands-On Machine Learning with Scikit-Learn & Tensorflow (Aurélien Géron)



Teaching method

Inverted classroom based


- Introduction lectures
- Real content you learn yourself with the notebooks. *Either to put in practice your knowledge or to learn ahead of another lecture*



Why

- Supposed to be better
- More fun
- Learning by doing

To give back sense to being present (Marcel Lebrun)



Tutorial I : Introduction to torch

[Link](#)


→ Copy to drive

Introduction

- **Goal** : introduction to pytorch
- **Program** : inverted classroom style
 - Theory
 - Overview to get the big picture of the Notebook
 - Work alone or in groups
- **Technical** : Google Colab, Pytorch


Theory

Gradient Descent Illustration




Training

- *Iterative* process






Learning curve



Tools


data structure & analysis



Seaborn

pandas

$$y_{it} = \beta' x_{it} + \mu_i + \epsilon_{it}$$



xarray



scikit
learn



scikit-image
image processing in python

main package for scientific computing in Python



NumPy



matplotlib



jupyter

interactive coding environments embedded in a webpage



IP[y]:
IPython

famous library to plot graphs in Python

Python-based ecosystem of open-source software for mathematics, science, and engineering.

h5py : common package to interact with a dataset that is stored on an H5 file

provides simple and efficient tools for data mining and data analysis

Overview of the notebook


Tutorial I (1)

- 1) Load necessary **libraries** (common libraries) and data
- 2) Create **model** (class **SimpleModel**)
 - $y = x * (x+2)$
- 3) **Run** the model
 - Test it, look at the output type
- 4) **Tensor operations**
 - New model : $\text{sum}_i(x_i+2)$ (**class SimpleModel2**)
 - Same for several 1D arrays at once (**class SimpleModel3**) : add axis=1
- 5) **Exercise 1** : calculate mean of array's elements

Tutorial I (1)

6) Optimization :

- New parabolic function to optimize (**class FLayer**) – trainable parameter is t initialized to 0
- *Gradient descent* applied, SGD optimizer defined
- Plot $f(t)$ results



7) Exercise 2 : change the parabolic function, the alpha learning rate,...

Exercise 1

```
# Define the MeanModel class
class MeanModel(nn.Module):
    def __init__(self):
        super(MeanModel, self).__init__()

    def forward(self, x):
        return torch.mean(x)

# Define data
arr = torch.tensor([[1, 2, 3, 4, 5], [2, 3, 4, 5.1, 6], [25, 65, 12, 12, 32], [90, 80, 70, 60, 50]])

# Instantiate the model
model = MeanModel()

# Run the model
result = model(arr)
```



 Copier le code

Tutorial II : Optimization and NN introduction

[Link](#)


→ Copy to drive

Introduction

- **Goal** : see how to do optimization in torch and NN introduction
- **Program** : inverted classroom style
 - Theory
 - Overview to get the big picture of the Notebook
 - Work alone or in groups
- **Technical** : Google Colab, Pytorch

Theory


Network model



Sigmoid and softmax

Sigmoid (*two-class* classifier) :

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



Softmax (*multi-class* classifier) :

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

Optimization

- Given a task we define

- Training data

$$\{x^i, y^i\}_{i=1,\dots,m}$$

- Network

$$f(x; \theta)$$

- Cost function


$$J(\theta) = \sum_{i=1}^m \text{loss}(y^i, f(x^i; \theta))$$

- Parameter initialization (weights, biases)

- random weights, biases initialized to small values (0.1)*

- Next, we *optimize the network parameters θ* (training)


- In addition, we have to set values for *hyperparameters*



Overview of the notebook

Tutorial II (1)

- 1) Load necessary **libraries** (common libraries)
- 2) **Linear regression**
 - Generate data points (80 training, 20 test)
 - Linear function (**class Linear**)
 - Loss function (**def loss_f**)
 - Train the model
 - Evaluate the model
- 3) **Exercise 1** : play with the parameters of the linear regression and the batch size



Tutorial II (2)

4) **Explanation** of the training loop with pseudocode

5) Building blocks of a Neural network

- Model (**class Dense**)
- Forward pass

5) Build a NN

- Multilayer NN (**dense1, dense2**)
- Overall model (**class MyModel**)

Tutorial III : Fully connected NNs

[Link](#)

→ Copy to drive

Introduction

- **Goal** : handwritten digit recognition with a fully connected NN
- **Program** : inverted classroom style
 - Theory
 - Overview to get the big picture of the Notebook
 - Work alone or in groups
- **Technical** : Google Colab, Pytorch


Theory

MNIST Dataset

- MNIST database with hand-written digits
- 60000 training images and 10000 testing images
- Created in 1994
- 128x128 binary images




Network model



Sigmoid and softmax

Sigmoid (*two-class* classifier) :

$$\sigma(x) = \frac{1}{1 + e^{-x}}$$



Softmax (*multi-class* classifier) :

$$\text{softmax}(z)_i = \frac{\exp(z_i)}{\sum_j \exp(z_j)}$$

Optimization

- Given a task we define

- Training data

$$\{x^i, y^i\}_{i=1,\dots,m}$$

- Network

$$f(x; \theta)$$

- Cost function


$$J(\theta) = \sum_{i=1}^m \text{loss}(y^i, f(x^i; \theta))$$

- Parameter initialization (weights, biases)

- random weights, biases initialized to small values (0.1)*

- Next, we *optimize the network parameters θ* (training)

- In addition, we have to set values for *hyperparameters*



Overview of the notebook

Tutorial III (1)

- 1) Load necessary **libraries** (common libraries) and data
- 2) **Training loop** (similar to Tutorial II)
 - Explanation
- 3) **Building blocks of a NN** (similar to Tutorial II)
- 4) **Structure of a NN**
 - Definition of a model and the forward pass (**class MyModel**)

Tutorial III (2)

5) Load the **data** (MNIST dataset)

- Training data, test data, normalization
- Plot examples

6) Build a NN

- Model, loss function, optimizer
- Training function
- Testing function
- Train the model (loss curves)
- Get the accuracy of the model

Tutorial III (2)

7) [Exercise 1](#) : build a NN with two layers

Exercises

```

import torch
import torch.nn as nn
import torch.optim as optim
from torch.utils.data import DataLoader, TensorDataset

# Définition du réseau
class SimpleNN(nn.Module):
    def __init__(self):
        super(SimpleNN, self).__init__()
        self.fc1 = nn.Linear(784, 1500) # Input size = 784 (ex: MNIST), output = 1500
        self.relu = nn.ReLU()
        self.fc2 = nn.Linear(1500, 10) # Output size = 10 (classes)
        self.softmax = nn.Softmax(dim=1)

    def forward(self, x):
        x = self.fc1(x)
        x = self.relu(x)
        x = self.fc2(x)
        x = self.softmax(x)
        return x

# Données fictives pour l'exemple
# 784 features (ex: pixels de MNIST), 1000 échantillons
X = torch.rand(1000, 784)
y = torch.randint(0, 10, (1000,)) # Classes entre 0 et 9

# Dataset et DataLoader
dataset = TensorDataset(X, y)
dataloader = DataLoader(dataset, batch_size=64, shuffle=True)

# Initialisation du modèle, de la fonction de perte et de l'optimiseur
model = SimpleNN()
criterion = nn.CrossEntropyLoss() # Perte adaptée pour classification
optimizer = optim.Adam(model.parameters(), lr=0.001) # Taux d'apprentissage initial

```

Exercise 1

```

# Entraînement du modèle
def train_model(dataloader, model, criterion, optimizer, epochs=10, lr_scheduler=None):
    for epoch in range(epochs):
        total_loss = 0
        for inputs, labels in dataloader:
            optimizer.zero_grad() # Réinitialiser les gradients
            outputs = model(inputs) # Passage avant
            loss = criterion(outputs, labels) # Calcul de la perte
            loss.backward() # Calcul des gradients
            optimizer.step() # Mise à jour des paramètres
            total_loss += loss.item()

    # Ajustement dynamique du taux d'apprentissage
    if lr_scheduler:
        lr_scheduler.step()

    print(f"Epoch {epoch+1}/{epochs}, Loss: {total_loss:.4f}, LR: {optimizer.param_group[0]['lr']}")

# Ajustement du Learning rate avec ReduceLROnPlateau
scheduler = optim.lr_scheduler.ReduceLROnPlateau(optimizer, mode='min', factor=0.5, patience=10)

# Entrainer le modèle
train_model(dataloader, model, criterion, optimizer, epochs=10, lr_scheduler=scheduler)

```

Tutorial IV : Convolutions

[Link](#)


→ Copy

Introduction


- **Goal** : basics to perform image recognition (Inception)
- **Program** : inverted classroom style
 - Theory
 - Overview to get the big picture of the Notebook
 - Work alone or in groups
- **Technical** : Google Colab, Pytorch

Theory


Human vision




Receptive fields size



Features




Computer vision




Kernel (filter)

- Used to **detect features** (vertical/horizontal filter,...)
- Different kernels to create different feature maps → learn to see various patterns and details in images

1	0	1
0	1	0
1	0	1



Image



Convolved
Feature

how well the pattern in the kernel matches the content in that part of the image

= Feature map


Kernel example

1	0	-1
1	0	-1
1	0	-1

Vertical

1	1	1
0	0	0
-1	-1	-1

Horizontal





Vertical edges




Horizontal edges

Padding, Stride, Dilation



Inception V3 model

- Deep learning model based on Convolutional Neural Networks
- Used for image classification
- Released in year 2015
- It has 42 layers








Overview of the notebook

Tutorial IV (1)


- 1) Load necessary **libraries** (common libraries and personal modules)
- 2) Images
- 3) **Convolutions**

Perceptron



Tutorial IV (2)

- Pre-processing of the image
 - Gray-scale, cropping, float conversion, normalization
 - Add dimensions (batch, channel, height, width) (in `get_convolved`)
 - Convert Numpy to pytorch (in `get_convolved`)
- Define the convolution
 - Forward model (`class Model`)
 - Apply 4 convolutions one after each other (inside `class Model`, call `conv_2d` function)
- Define the filter
 - Identity filter
 - Convert to `np.array` (in `get_convolved`)
 - Add dimensions
 - Convert Numpy to pytorch



```
flt_mtx = [
    [ 0, 0, 0, 0, 0, ],
    [ 0, 0, 0, 0, 0, ],
    [ 0, 0, 1, 0, 0, ],
    [ 0, 0, 0, 0, 0, ],
    [ 0, 0, 0, 0, 0, ],
] # identity transformation
```

Tutorial IV (3)

- Use it ! (`ims_convolved = get_convolved(img_raw, flt_mtx)`)

- You get back 5 figures

• Exercise :

1. experiment with different filters and understand what they do, e.g.:

- identity transformation
- identity transformation with positive non-unit values
- identity transformation with negative unit value
- identity transformation off center
- blurring with box filter
- edge detection with + and - bands
- try whatever you like

2. experiment with convolution parameters:

- padding = 1, 2, 3
- stride = 2
- dilation = 2

0	0	0	0	0
0	0	0	0	0
0	0	30	0	0
0	0	0	0	0
0	0	0	0	0

0	0	0	0	0
0	0	0	0	0
0	0	-1	0	0
0	0	0	0	0
0	0	0	0	0

0	0	0	0	0
0	0	0	0	0
0	0	0	0	0
0	0	1	0	0
0	0	0	0	0

1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1
1	1	1	1	1

0	-1	1	0	0
0	-1	1	0	0
0	-1	1	0	0
0	-1	1	0	0
0	-1	1	0	0

Tutorial IV (4)

- **Most common filters**

- Define 1D functions
- Create 2D filters by repeating the 1D filters along axis 0 (`np.tile`)
- Multiply by `transpose()` to get the horizontal dimension (filter size does not change)
- Use them ! (`ims_convolved = get_convolved(img_raw, flt_mtx)`)

4) Homework (leave it for now)

filter type	effect
gaussian	bluring
first derivative of gaussian	detection of edges
second derivative of gaussian	detection of peaks

Tutorial IV (5)

- 5) Load a **pretrained model** (**inception V3**) from **torchhub**
- 6) **Test** the model
 - Preprocessing of the image (cropping, shuffle sizes, totensor, normalize adds batch size) → [1,3,299,299]
 - Desactivate the gradient (eval mode)
 - Get the logits (**1000** values), then the probabilities (applying softmax)
 - Print out the 5 most probable classes
 - Do the same with 100 classes

Tutorial V : Transfer Learning

[Link](#)


Introduction

- **Goal** : use the Inception model to classify images of different nature (it/de), learn how to save a model
- **Program** : inverted classroom style
 - Theory
 - Overview to get the big picture of the notebook
 - Work alone or in groups
- **Technical** : Google Colab, Pytorch

Theory


Transfer Learning

- Try to find an existing neural network that accomplishes a **similar task** to the one you are trying to tackle
- reuse the lower layers of this network
 - Output layer should usually be replaced
- **Speeds up** training and requires much fewer training data




Training Loop

- *Iterative* process




Learning curve



*Measure the error contribution
from each connection*

Dataset Augmentation

- Apply **realistic transformations** to data to create new synthetic samples, with same label



Overview of the notebook


Tutorial V (1)

- 1) Load necessary **libraries** (common libraries and personal modules)
- 2) **Transfer Learning**
 - Load the **inception model** (**base_model**)
 - Build **new model** using the **base_model** (**model**)
 - Define a head function (`in_features, n_classes=2`) : use of sigmoid
 - **Optimization** : define the loss function (**criterion**) and the optimizer
 - **Helper functions** to get the prediction (**get_predictions**) and to compute the batch accuracy (**calculate_accuracy_batch**)

Tutorial V (2)

3) Dataset

- **Images :**
 - Get them from `ML3` folder
 - Preprocessing using `transforms.Compose (resize, tensor, normalize)`
 - Shape `[3,299,299]`
 - Transform to numpy array for display (`im_numpy`)
- **Labels**
- **Split** dataset into train/test samples
 - `Torch.utils.data.random_split`
 - `Train_test_split` with `stratify` enabled from scikit-learn
- Create **data loaders** for training/val (beware shuffle param)
- Example : batch of 10 images, plot them, print logits and output classes (`res`)



Tutorial V (3)

4) Training

- `Train_model` function
 - Contains the loop on `epochs`
 - Calls the `train` and `validate` functions (beware the params)
 - Save `history` (loss, accuracy) and model for a given epoch
 - `Train function` : reset gradients, compute logits and loss, compute gradients (backward prop), update parameters, calculate accuracy of the batch (helper functions), returns train loss and train accuracy
 - `Validate function`: structure BUT (no optimizer, no backward, no update of the params), returns test loss and test accuracy
- `Plot_history` function
- Run it ! `history = train_model(...)` using 70 epochs

Tutorial V (4)

5) Load trained variables from checkpoint

- Choose epochs values
- Load corresponding models
- Call validate function to get the validation loss and validation accuracy (see how it evolves)

6) Save final model for inference

7) Inference :

- Load the model, eval() mode
- Get an image, preprocess (convert to tensor, add batch dimension)
- Get the logits and associated class

```
#loaded_model =  
torch.load('inference_model.pth')  
#loaded_model =  
loaded_model.to(device)  
#loaded_model.eval(); # set the  
model to inference mode
```

```
loaded_model = torch.load(  
    "inference_model.pth",  
    map_location=device,  
    weights_only=False  
)  
loaded_model.eval()
```

Tutorial V (4)

8) Improve the results : data augmentation

- Load images from `ML3` folder
- Preprocess (`resize`, `Randomcrop`, `tensor`, `normalize`)
- Convert to Numpy for plotting purposes (`im_numpy`)
- The rest of the code is similar to previous code

9) Exercise

Tutorial VI : RNNs

[Link](#)

Introduction

- **Goal** : use Recurrent Neural Networks to predict and generate text sequence
- **Program** : inverted classroom style
 - Theory
 - Overview to get the big picture of the notebook
 - Work alone or in groups
- **Technical** : Google Colab, Pytorch

Theory

Overview of the notebook

Tutorial VI (1)

1) Load necessary **libraries** (common libraries and personal modules)

2) **Text data**

- Read the data (`rnn.txt`), print the first 100 words

3) **Build the dataset**

- 2 dictionaries : word → id (`dictionary`) and id → word (`reverse_dictionary`)
 - `Build_dictionaries` function
- Vocabulary size = 493 (0=most common word)
- Helper functions to get sequences of int or words (`text_to_ints`, `ints_to_text`)
- Print example : first 100 words (or int), length of input data=2118 (`words_as_int`)

Tutorial VI (2)

3) Data streaming

- Create dataset using `WordDataSet` class to create blocks of text
 - Block length = $n_input+1 = 3+1 = 4$
- Create `DataLoader` with `batch size=50`, `preprocess` data (separate input and target sequences, `stack` data and convert Numpy \rightarrow `tensors`, put them to `GPU`)
 - Length of dataset = Number of blocks in sequence = Total length/Block length = $2118/4=529$
- Example : print the 50 samples that are in the 1st batch

Tutorial VI (3)

4) Construct model

- Create class RNN
 - Embedding layer (`vocab_size`, `embedding_dim=128`)
 - Loop to add the 3 LSTM layers
 - FC layer with `vocab_size`
- Define sequence of 3 words (`n_input`), define dimension of 3 LSTM, call the RNN class
- Investigate the model using Tensorboard
 - Create an input of size=5 (`x`), transform to tensor, add batch
 - `SummaryWriter` to save the model and be able to open it with tensorboard
 - Print the output size of `y` : [5,1,493]
- Test the NOT trained model and see that it is bad
 - Get the first batch (break), apply the model, get the predictions, compare to true

Tutorial VI (4)

5) Train the model

- Params : `n_input = 3, batch_size=50 , one LSTM layer (128),n_epochs=200`
- Create the `data loader` (preprocess data)
- Optimization part : `criterion, optimizer` (RMSPprop)
- Training loop on epochs, then on batches
 - Initialize gradients
 - Get the output (`seq_len, batch_size, vocab_size`), reshape it to `(seq_len*batch_size, vocab_size)`
 - Reshape labels (`seq_len, batch_size`) to `(seq_len*batch_size)`
 - Compute the loss between output and true labels
 - Backward prop, param update
 - Compute loss and accuracy
- Plot loss and accuracy

Tutorial VI (5)

6) Generate text with RNN

- Function to generate text (gen_long)
 - Input parameters : model, input sequence, number of words to generate (128)
 - No_grad() because we are in an evaluation mode
 - Loop on number of words, convert to tensor, predict, ...)